

Real-Time Systems¹

Sanjit Kumar Roy

Dept. of Information Technology

UPPSALA
UNIVERSITET

December 11, 2024



¹Course ID: 1DT004 11210 and I 1DT063 11202 (HT2024)

Multiprocessor Scheduling - Day 2



UPPSALA
UNIVERSITET

- Multiprocessor scheduling
 - Global scheduling
 - Partitioned scheduling
 - Semi-partitioned scheduling




Multiprocessor Scheduling



- Global scheduling (**full migration**)
 - A task is allowed to execute on an arbitrary processor (sometimes even after being preempted)
- Partitioned scheduling (**no migration**)
 - Each instance of a task must execute on the same processor
 - Equivalent to multiple uniprocessor systems!
- Semi-partitioned scheduling
 - Static task assignment
 - Each instance (or part of it) of a task is assigned to a fixed processor
 - Task instance or part of it may migrate






Multiprocessor Scheduling contd.




Multiprocessor Scheduling contd.



Multiprocessor Scheduling contd.



Global Scheduling



- All ready tasks are kept in a global queue
- When selected for execution, a task can be assigned to any processor, even after being preempted (task migration to another processor!)
 - Task migration may add overheads

Global Scheduling contd.

Any algorithm for single processor scheduling may work, but schedulability analysis is non-trivial

- EDF (optimal for single processor)
 - Unfortunately **EDF is not optimal for multiprocessors**
 - Example: $\tau = \{(C_i, T_i)\} = \{(6, 12), (6, 12), (8, 13)\}$; Number of processors $m = 2$
 - No exact, but only sufficient schedulability test known
- Fixed Priority Scheduling (e.g. RM)
 - Difficult to find the optimal priority order
 - Difficult to check the schedulability
 - Suffer from the **Dhall's effect**



Dhall's effect (Dhall & Liu, 1978):

- With RM, DM and EDF, some low-utilization task sets can be unschedulable regardless of how many processors are used
 - Example: $\tau = \{(C_i, T_i)\} = \{(2\epsilon, 1), (2\epsilon, 1), (2\epsilon, 1), (1, 1 + \epsilon)\}$
 - RM, DM, EDF can't schedule τ on any number of processors



Dhall's effect (Dhall & Liu, 1978):

- With RM, DM and EDF, some low-utilization task sets can be unschedulable regardless of how many processors are used
 - Example: $\tau = \{(C_i, T_i)\} = \{(2\epsilon, 1), (2\epsilon, 1), (2\epsilon, 1), (1, 1 + \epsilon)\}$
 - RM, DM, EDF can't schedule τ on any number of processors
- Utilization:
 - Number of processors, m
 - A set of $m + 1$ tasks
 - m tasks each having utilization $2\epsilon/1$
 - One task has utilization $1/(1 + \epsilon)$



Dhall's effect (Dhall & Liu, 1978):

- With RM, DM and EDF, some low-utilization task sets can be unschedulable regardless of how many processors are used
 - Example: $\tau = \{(C_i, T_i)\} = \{(2\epsilon, 1), (2\epsilon, 1), (2\epsilon, 1), (1, 1 + \epsilon)\}$
 - RM, DM, EDF can't schedule τ on any number of processors
- Utilization:
 - Number of processors, m
 - A set of $m + 1$ tasks
 - m tasks each having utilization $2\epsilon/1$
 - One task has utilization $1/(1 + \epsilon)$

$$U = \frac{m \times 2\epsilon + \frac{1}{1+\epsilon}}{m} \rightarrow 0 \quad (\text{when } \epsilon \rightarrow 0 \text{ and } m \rightarrow \infty)$$

Global Scheduling contd.

How can we avoid Dhall's effect:

- RM, DM & EDF only account for task deadlines! Actual computation demands are not accounted for
- Dhall's effect can easily be avoided by letting tasks with high utilization receive higher priority:



Global Scheduling contd.

How can we avoid Dhall's effect:


- RM, DM & EDF only account for task deadlines! Actual computation demands are not accounted for
- Dhall's effect can easily be avoided by letting tasks with high utilization receive higher priority:

New priority assignment scheme: (Andersson, Baruah & Jonsson, 2001)

$RM - US\{m/(3m - 2)\}$

- $RM - US\{m/(3m - 2)\}$ assign (static) priorities to tasks according to the following rule:
 - τ_i has **highest priority** if $U_i > m/(3m - 2)$
 - τ_i has **RM priority** if $U_i \leq m/(3m - 2)$
 - Here, $U_i = C_i/T_i$, and m is the number of processor

Partitioned Scheduling



- Two steps:
 - Determine a mapping of tasks to processors
 - Perform run-time scheduling

Partitioned Scheduling contd.

Bin-packing Algorithms:

- Pack items of different sizes into a minimum number of bins, each with a fixed capacity
 - Solutions (Heuristics): First Fit, Next Fit, Best Fit
- Application to multiprocessor systems:
 - Bins are represented by processors and objects by tasks
 - The decision whether a processor is “fully loaded” or not is derived from a utilization-based schedulability test



Partitioned Scheduling contd.

Partitioned with EDF:

- Assign tasks to processors such that processor capacity does not exceeded (i.e. utilization bounded by 1)
- Schedule each processor using EDF



Partitioned Scheduling contd.

Partitioned with RM:

Rate-Monotonic First Fit (RMFF): (Dhall & Liu, 1978)

- First sort all tasks in non-decreasing order of their periods
- Task Assignment
 - Start from the highest priority task, and assign tasks to processors in the First Fit manner
 - A task τ_i can be assigned to a processor P_k if all the tasks assigned to P_k are RM-schedulable i.e.,
 - Total utilization of tasks assigned on P_k is bounded by $n(2^{1/n} - 1)$, where n is the number of tasks assigned
 - One may also use the precise test (response time analysis) to get a better assignment!
 - Add a new processor if needed for the RM-test



UPPSALA
UNIVERSITET

Partitioned Scheduling contd.

Utilization bounded by 50%

- The Partitioning Problem is similar to Bin-packing problem (NP-hard)
- Limited Resource Usage, 50%
- Example:
 - Number of processors, m
 - A set of $m + 1$ tasks each having utilization $0.5 + \epsilon$




Partitioned Scheduling contd.

Utilization bounded by 50%

- The Partitioning Problem is similar to Bin-packing problem (NP-hard)
- Limited Resource Usage, 50%
- Example:
 - Number of processors, m
 - A set of $m + 1$ tasks each having utilization $0.5 + \epsilon$

$$U = \frac{(m + 1) \times (0.5 + \epsilon)}{m} \rightarrow 0.5 \quad (\text{when } \epsilon \rightarrow 0 \text{ and } m \rightarrow \infty)$$

Semi-partitioned Scheduling



- Bin-packing with item splitting
- Resource can be “fully” (better) utilized

Semi-partitioned Scheduling contd.

Lakshmanan's Algorithm [ECRTS'09]:

- Sort all tasks in **decreasing order of their utilization**
- Pick up one processor, and assign as many tasks as possible from high utilization task to low utilization task
- If a task does not fit in the current processor then split the task into two
 - Assign one part of the task to the current processor
 - Pickup an empty processor, and assign another part of the task into it
- Utilization bound: **65%**



Semi-partitioned Scheduling contd.

Breadth-First Partitioning Algorithms [RTAS 2010]

- Sort all tasks in **increasing order of their priority**
- Select the processor on which the assigned utilization is the lowest, and assign tasks from low priority to high priority
- If a task does not fit in the selected processor then split the task
- Utilization bound: **69%**

